

Minh Duc (David) Chu

davidchu11381@gmail.com

 [LinkedIn](#)  [GitHub](#)  [Google Scholar](#)

Research Interests

AI Safety and Alignment, Conversational AI, Societal Impacts of LLMs, Computational Social Science

Education

University of Southern California, Los Angeles, California Aug 2023 - Present
Ph.D. Candidate in Computer Science • Cumulative GPA: 3.95
Advisors: Luca Luceri, Kristina Lerman • Expected Graduation: Dec 2027
Carleton College, Northfield, Minnesota Sep 2019 - Jun 2023
B.A. with Distinction in Computer Science & Statistics • Cumulative GPA: 3.75

Work Experience

Character.AI, Redwood City, California Start Date: Sep 2026
Incoming Research Engineer Intern, AI Safety and Alignment Team
Mentors: Deniz Demir, Dylan Hadfield-Menell
Projects: TBD
Anthropic, Berkeley, California Start Date: May 2026
Incoming AI Safety Fellow
Mentors: TBD
Projects: TBD

Selected Publications & Preprints

Illusion of Intimacy: How Emotional Dynamics Shape Human-AI Relationships
Minh Duc Chu, Patrick Gerard, Kshitij Pawar, Charles Bickham, Kristina Lerman. In preparation for **ACII 2026** [\[Preprint\]](#) [\[Media Coverage\]](#)
Developed emotional mirroring and synchrony metrics over 17,822 emotionally salient companion-chatbot dialogues, revealing strong affect tracking and frequent failure to de-escalate explicit and self-harm language
BigTokDetect: A Clinically Informed Vision-Language Modeling Framework for Detecting Pro-Bigorexia Videos on TikTok
Minh Duc Chu, Kshitij Pawar, Zihao He, Roxanna Sharifi, Ross Sonnenblick, Magdalayna Curry, Laura D'Adamo, Lindsay Young, Stuart B Murray, Kristina Lerman. In Proceedings of **EACL'26** [\[Preprint\]](#)
Led a multimodal benchmark and clinically grounded harm-taxonomy that operationalizes psychosocial risk in short-form videos by modeling the joint signal from visuals and text, then benchmarking and domain-fine-tuning VLMs to enable scalable detection and moderation of pro-bigorexia (muscle dysmorphia) content
Improving and Assessing the Fidelity of LLMs Alignment to Online Communities
Minh Duc Chu, Zihao He, Rebecca Dorn, Kristina Lerman. In Proceedings of **NAACL'25** [\[Paper\]](#) [\[Code\]](#)
Led an instruction-tuning and evaluation pipeline that aligns LLMs to community-specific language, then rigorously audits fidelity and robustness across various high-level dimensions, including psychometric testing across aligned LLMs

Services

PC/Reviewer: NAACL, ACL, EMNLP, AACL, EACL, CSCW, WWW, ICWSM, ICWSM'26 (**Local Chair**)
Teaching: DSCI-352: Applied Machine Learning and Data Mining (Fall 2025), DSCI-549: Introduction to Computational Thinking and Data Science (Spring 2025), DSCI-550: Data Science at Scale (Spring 2026)

Skills

Programming: Python, C++, C, JavaScript, Java, HTML, Docker, CUDA
Statistics: Bayesian Statistics, Probability, Time Series, Spatial Statistics, Visualization, Sampling Techniques
Machine Learning: Python (scikit-learn, PyTorch, TensorFlow, Keras, OpenCV, Huggingface), R (caret, bayesrules, TSA)
LLM: Preference Optimization (GRPO, DPO), RLHF, RAG, AutoGen, Multi-Agent, LLM-as-a-Judge, Alignment, Safety, Distributed Training, Multimodal Large Models